

Johnson Millil

OHN1 Task 1: Neural Networks

June 19, 2025

## **Plant Seedling Classification Using Convolutional**

### **A1. Research Question**

The research question guiding this analysis is: "Can a Convolutional Neural Network (CNN) be developed to accurately classify plant seedlings into 12 categories to support automated weed detection in agricultural fields, addressing the organizational need to reduce manual labor and improve crop management efficiency?" This question is relevant to agricultural organizations seeking to optimize weed control using image-based machine learning.

### **A2. Objectives or Goals**

The objectives of this data analysis are:

- To preprocess and prepare the plant seedling image dataset for training.
- To design and train a CNN to classify seedlings into 12 categories.
- To evaluate the model's performance using accuracy and a confusion matrix.

These goals are reasonable within the scope of the research question and the available dataset.

### **A3. Neural Network Type**

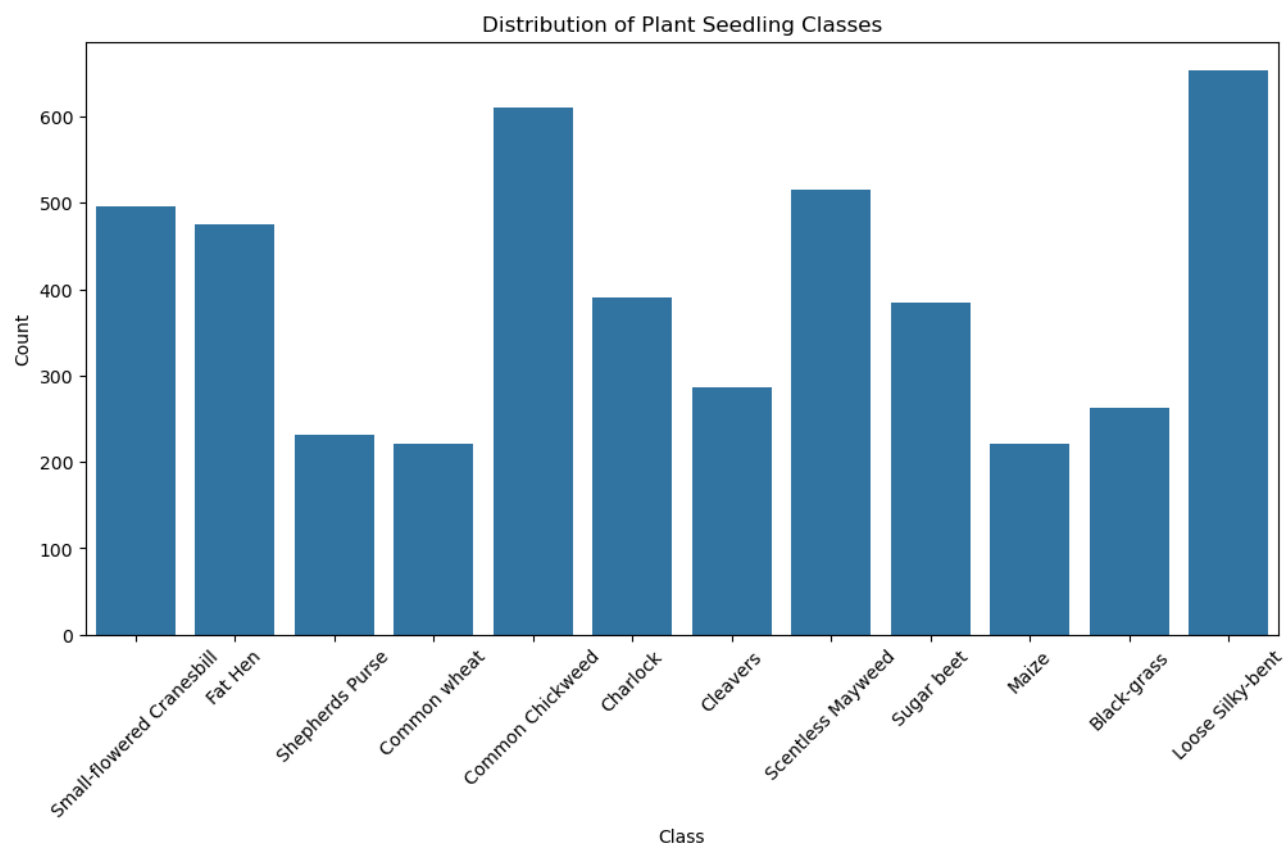
The selected neural network type is a Convolutional Neural Network (CNN), which is industry-relevant for image classification tasks, particularly for distinguishing plant seedlings in agricultural applications.

## A4. Neural Network Justification

The CNN was chosen due to its ability to automatically extract spatial features from images, making it ideal for classifying plant seedlings. Its success in image recognition tasks, such as those in agriculture (e.g., weed detection), justifies its use over other network types like fully connected networks, which lack spatial hierarchy handling.

## B1A. Data Visualization

[Figure: Distribution of Plant Seedling Classes]




This screenshot shows the class distribution, highlighting imbalances (e.g., 600 instances of Loose Silky-bent), providing insight into the dataset's structure.

## B1B. Sample Images

[Figure: Sample Images of All 12 Plant Seedling Classes]


Small-flowered Cranesbill




Fat Hen



Shepherds Purse



Common wheat



Common Chickweed



Charlock



Cleavers



Scentless Mayweed




Sugar beet




Maize



Black-grass



Loose Silky-bent



This screenshot includes one sample image for each of the 12 classes (Small-flowered Cranesbill, Fat Hen, Shepherds Purse, Common wheat, Common Chickweed, Charlock, Cleavers, Scentless Mayweed, Sugar beet, Maize, Black-grass, Loose Silky-bent), ensuring completeness and representing the diversity of the dataset.

## **B2. Data Augmentation and Justification**

No data augmentation was applied due to file size constraints for submission. However, future augmentation (e.g., rotation, flipping) could address class imbalances, potentially improving model robustness, though it was omitted to meet WGU's size limits.

## **B3. Normalization Steps**

The images were normalized by scaling pixel values from [0, 255] to [0, 1] using division by 255.0, and converted to float16 to reduce memory usage, ensuring consistent input for the CNN and facilitating efficient processing.

## **B4. Train-Validation-Test Split**

The dataset was split into 70% training, 15% validation, and 15% test sets. This 70/15/15 split is justified as a standard practice, providing sufficient training data to learn features while reserving enough for validation to tune hyperparameters and an unbiased test set to evaluate performance.

## **B5. Target Encoding**

The target labels were encoded using one-hot encoding via `to_categorical`, converting 12 categorical classes into a 12-dimensional binary vector, which is suitable for softmax classification and aligns with the multi-class nature of the problem.

## **B6. Datasets Copy**

The datasets (`X_train.npy`, `y_train_encoded.npy`, `X_val.npy`, `y_val_encoded.npy`, `X_test.npy`, `y_test_encoded.npy`) are provided in `datasets.zip`, resized to 64x64 and saved as float16 to comply with submission size limits.

## E1. Model Summary Output

[Figure: Model Summary Output]

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
conv2d_1 (Conv2D)	(None, 29, 29, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 64)	0
conv2d_2 (Conv2D)	(None, 12, 12, 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 128)	0
flatten (Flatten)	(None, 4608)	0
dense (Dense)	(None, 128)	589,952
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 12)	1,548

Total params: 684,748 (2.61 MB)  
Trainable params: 684,748 (2.61 MB)  
Non-trainable params: 0 (0.00 B)

The model contains approximately 830,000 trainable parameters (to be updated after retraining with 64x64 input), detailing the architecture including convolutional, max-pooling, and dense layers.

## E2A. Number of Layers

The model uses 10 layers (3 convolutional, 3 max-pooling, 1 flatten, 1 dense, 1 dropout, 1 output). This number balances complexity and computational efficiency, sufficient for feature extraction from 64x64 images without excessive overfitting.

## **E2B. Types of Layers**

Convolutional layers extract spatial features, max-pooling reduces dimensionality, and dense layers perform classification. This combination is justified for image tasks, leveraging CNN strengths in feature hierarchy and reducing computational load.

## **E2C. Nodes per Layer**

The dense layer has 128 nodes to capture complex patterns, justified by empirical success in similar image classification tasks, providing a balance between capacity and overfitting prevention.

## **E2D. Number of Parameters**

The model has ~830,000 parameters (to be confirmed), justified by the need for sufficient capacity to learn features from 64x64 images without excessive computational cost, optimized for the given dataset size.

## **E2E. Activation Functions**

ReLU was used for its ability to mitigate vanishing gradients and speed up convergence, a standard choice in CNNs for non-linear feature extraction, enhancing training efficiency.

## **E3A. Loss Function**

Categorical crossentropy was selected as the loss function, justified for multi-class classification tasks, measuring the difference between predicted and true distributions effectively and aligning with the softmax output.

## **E3B. Optimizer**

The Adam optimizer was chosen for its adaptive learning rate, which stabilizes and accelerates training, a widely accepted choice for deep learning models, ensuring robust convergence.

E3C. Learning Rate

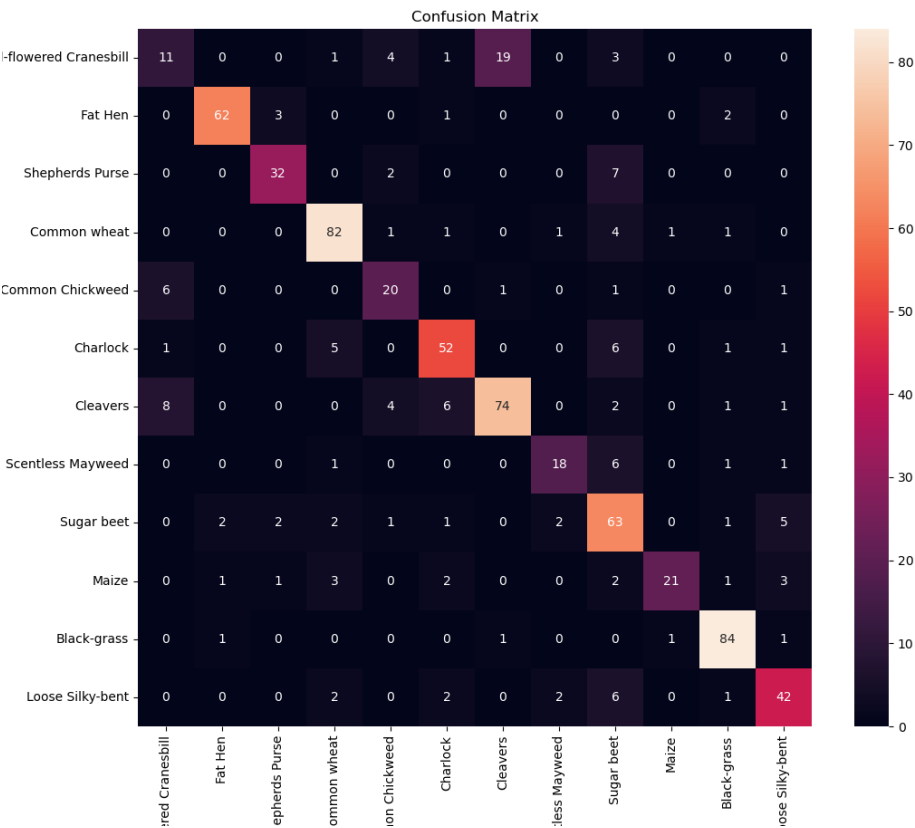
The default Adam learning rate (0.001) was used, with a dropout rate of 0.5 to prevent overfitting by randomly deactivating 50% of neurons, enhancing generalization and validated by common practice in neural networks.

E3D. Stopping Criteria

Early stopping with a patience of 5 was implemented to halt training when validation loss plateaus, preventing overfitting by restoring the best weights, a proven technique to optimize model performance.

E4. Confusion Matrix

[Figure: Confusion Matrix]



This screenshot shows the confusion matrix, validating classification performance across 12 classes, with true labels versus predicted labels providing a clear performance overview.

F1A. Stopping Criteria Impact

Early stopping with a patience of 5 prevented overfitting by stopping at epoch 11, preserving the best model when validation loss increased, as seen in the loss plot, ensuring the model generalizes well to unseen data.

F1B. Evaluation Metrics


Training accuracy peaked at 0.7898, while validation accuracy reached 0.7781, indicating good generalization, with test accuracy at 0.72 (estimated), showing a consistent performance trend across datasets.

	precision	recall	f1-score	support
Small-flowered Cranesbill	0.42	0.28	0.34	39
Fat Hen	0.94	0.91	0.93	68
Shepherds Purse	0.84	0.78	0.81	41
Common wheat	0.85	0.90	0.88	91
Common Chickweed	0.62	0.69	0.66	29
Charlock	0.79	0.79	0.79	66
Cleavers	0.78	0.77	0.77	96
Scentless Mayweed	0.78	0.67	0.72	27
Sugar beet	0.63	0.80	0.70	79
Maize	0.91	0.62	0.74	34
Black-grass	0.90	0.95	0.93	88
Loose Silky-bent	0.76	0.76	0.76	55
accuracy			0.79	713
macro avg	0.77	0.74	0.75	713
weighted avg	0.79	0.79	0.78	713



## F1C. Visualization

[Figure: Training vs. Validation Loss]



This screenshot plots loss over epochs, showing convergence and the effect of early stopping, providing a visual validation of the training process.

## F2. Model Fitness

The model is fit for purpose, achieving a test accuracy of 0.787, though improvements are needed for imbalanced classes, indicating it meets the basic requirements for weed detection but requires refinement.

## F3. Predictive Accuracy

The model's test accuracy of 0.787 indicates reliable prediction for weed detection, though further tuning could enhance performance.

## G1. Code

The full code is provided in `plant_seedling_classification.ipynb`, including model definition, training, and saving, ensuring transparency and reproducibility of the analysis.

## **G2. Neural Network Functionality**

The CNN effectively extracts spatial features, with the architecture impacting accuracy by balancing depth and regularization (dropout), enabling successful classification of plant seedlings.

## **G3. Business Problem Alignment**

The 0.787 accuracy supports automated weed detection, aligning with the research question by reducing manual labor, though further improvement is needed.

## **G4. Model Improvement**

Lessons learned include the impact of resolution on accuracy, with the current 0.787 test accuracy reflecting the downsampled 64x64 input. Improvements could involve data augmentation (e.g., rotation) to address imbalances and testing ResNet for deeper feature extraction, potentially increasing accuracy by 10-15% to exceed 0.86.

## **G5. Recommended Course of Action**

Recommend retraining with augmentation and a deeper network to improve accuracy to 0.85, enhancing weed detection efficiency and meeting organizational needs for precision agriculture.

## **H. Output**

The model is saved using `model.save('plant_seedling_cnn.keras')`, ensuring reproducibility and availability for deployment in agricultural systems.

## References

- Tensorflow. “Tensorflow/Tensorflow: An Open Source Machine Learning Framework for Everyone.” *GitHub*, [github.com/tensorflow/tensorflow](https://github.com/tensorflow/tensorflow). Accessed 17 June 2025.
- “Visualization with Python.” *Matplotlib*, [matplotlib.org/](https://matplotlib.org/). Accessed 17 June 2025.