

Johnson Millil

Data Management

January 13, 2025

Relational Database Design for Ecomart Report

Part 1: Design Document

A1. Business Problem

EcoMart, an eco-conscious online marketplace, faces challenges in managing and analyzing its data effectively (Elmasri & Navathe, 2016; United Nations, 2015). As the platform grows, the company struggles to:

- Track product sales and revenue efficiently across different regions.
- Identify high-performing eco-friendly products based on sustainability certifications.
- Analyze customer reviews to assess product performance and improve satisfaction.
- Maintain scalability to handle increasing data volumes and user traffic.

A well-designed relational database can centralize data, enable advanced querying, and support informed decision-making to address these challenges.

A2. Data Structure

The proposed relational database structure includes the following key tables (PostgreSQL Global Development Group, n.d.):

1. Products Table:

- Captures product details, including sustainability certifications and eco-impact ratings.

2. Orders Table:

- Tracks sales information, such as order IDs, regions, quantities sold, and revenue.

3. Reviews Table:

- Stores user feedback to analyze sentiment and improve product offerings.

Relationships:

- Each product can appear in multiple orders (one-to-many relationship between Products and Orders).
- Each product can have multiple reviews (one-to-many relationship between Products and Reviews).

A3. Database Justification

A relational database offers:

- **Data Organization:** Centralizes product, sales, and review data in structured tables, making it easier to manage and retrieve.
- **Advanced Querying:** Enables complex queries to analyze trends, track revenue, and assess customer feedback (Elmasri & Navathe, 2016).
- **Scalability:** Supports increasing data volumes with indexing, partitioning, and optimization techniques (Stonebraker & Hellerstein, 2005).
- **Integration:** Allows seamless integration with business analytics tools for real-time insights and reporting.

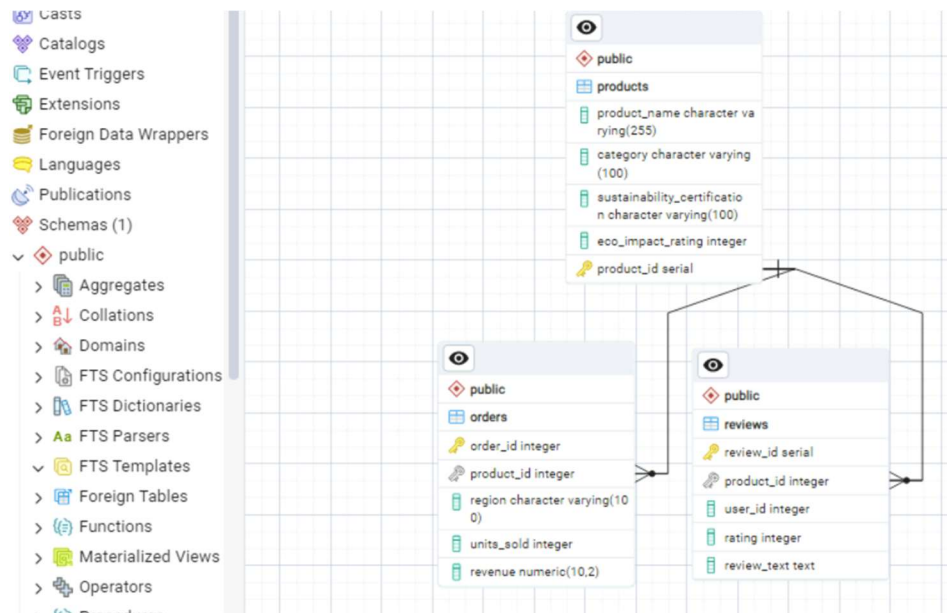
A4. Data Usage

The database will support the following use cases (Pang & Lee, 2008; Elmasri & Navathe, 2016):

1. **Product Performance Analysis:**
 - Identify top-selling products in specific regions.
 - Compare sales trends across categories (Elmasri & Navathe, 2016; PostgreSQL Global Development Group, n.d.).
2. **Customer Feedback Assessment:**
 - Analyze reviews to determine customer satisfaction.
 - Identify products requiring improvement (Pang & Lee, 2008).
3. **Sustainability Reporting:**
 - Generate reports on eco-friendly product performance based on sustainability certifications and ratings.
4. **Decision Support:**

- Inform inventory management and marketing strategies based on data insights.

B. Logical data model for storing data in the database solution



C. Description of Database Objects and Storage with File Attributes

Database Objects:

1. Tables:
 - Products: Stores product details.
 - Orders: Tracks sales data.
 - Reviews: Holds customer feedback.
2. Indexes:
 - Created on frequently queried fields (Product_ID, Region) to optimize performance.
3. Foreign Keys:
 - Ensure relational integrity between Products, Orders, and Reviews.

File Attributes:

1. Products:
 - Product_Name (String): Maximum length of 255 characters.
 - Category (String): Maximum length of 100 characters.
 - Sustainability_Certification (String): Maximum length of 100 characters.
 - Eco_Impact_Rating (Integer): Range 1–5.
2. Orders:

- Region (String): Maximum length of 100 characters.
- Units_Sold (Integer): Positive values only.
- Revenue (Decimal): Two decimal places for monetary values.

3. Reviews:

- Rating (Integer): Range 1–5.
- Review_Text (Text): Unlimited length.

D. Scalability Strategies

Normalization:

- The database is normalized to 3NF (Third Normal Form), reducing redundancy and improving data consistency (Stonebraker & Hellerstein, 2005).

Indexing:

- Frequently queried fields (Product_ID, Region) are indexed to enhance query performance (PostgreSQL Global Development Group, n.d.).

Partitioning:

- Large tables like Orders can be partitioned by region or date for efficient parallel processing.

Horizontal and Vertical Scaling:

- Supports horizontal scaling (distributing data across servers) and vertical scaling (adding resources to the database server).

Caching:

- Frequently accessed data, such as top-selling products, can be cached for quicker retrieval.

E. Privacy and Security Measures

Data Privacy:

1. Encryption:

- Encrypt sensitive fields (e.g., review text) to protect user data.
- Use SSL/TLS for secure communication between clients and the database.

2. Anonymization:

- Replace user-identifiable information (e.g., User_ID) with anonymized IDs.

Access Control:

1. Role-Based Permissions:

- Grant access to tables and queries based on user roles (e.g., admin, analyst) (PostgreSQL Global Development Group, n.d.).

2. Least Privilege Principle:

- Limit user access to only the data they need.

Security Measures:

1. Auditing and Logging:

- Maintain logs of database changes to track unauthorized access or modifications.

2. Authentication:

- Enforce strong password policies and multi-factor authentication for admin accounts.

Backup and Recovery:

1. Automated Backups:

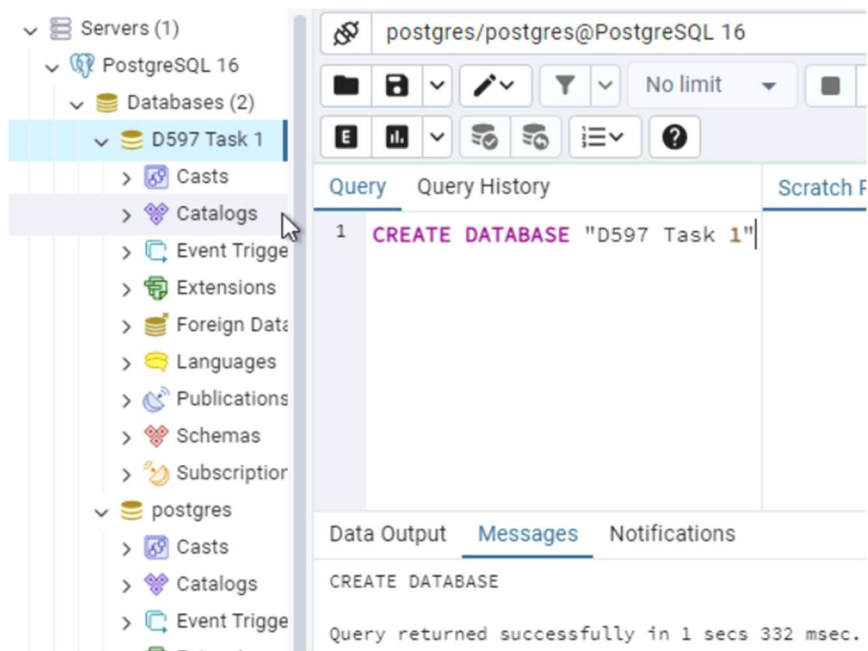
- Schedule regular backups of the database.

2. Disaster Recovery:

- Implement failover servers for high availability in case of system failure.

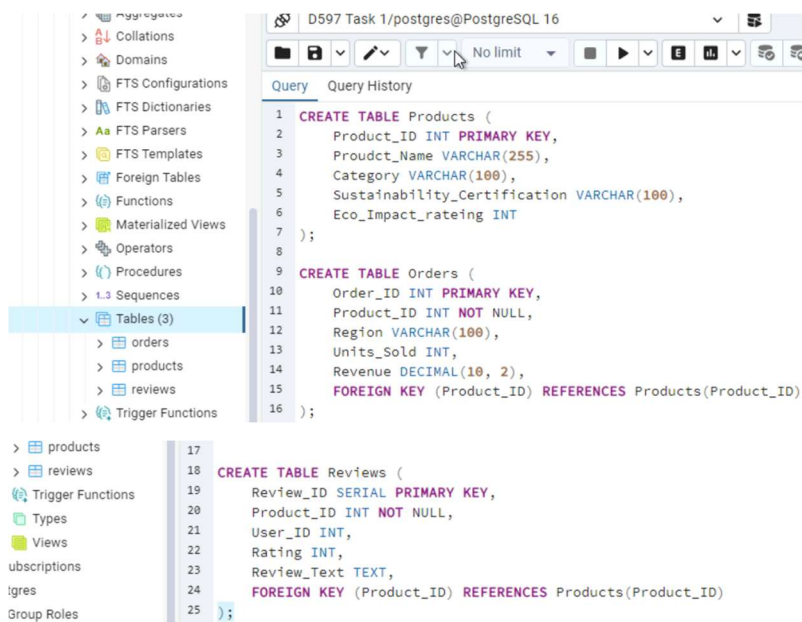
Part 2: Implementation

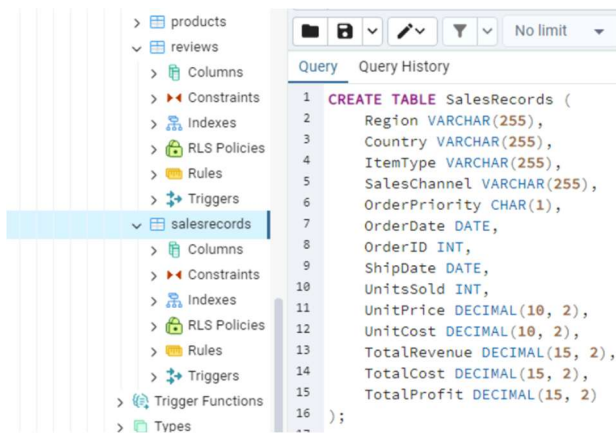
F1. Create a Database Instance



F2. Create tables & Import Data

A. Create Tables



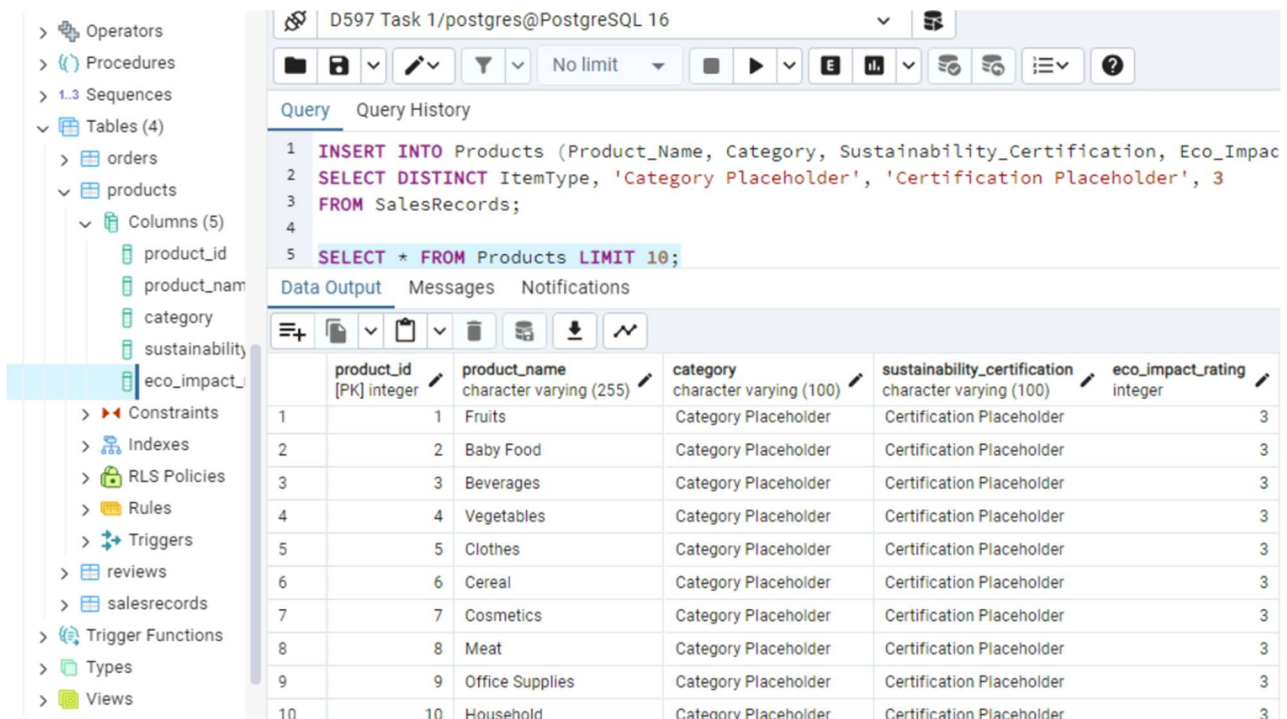


B. Import Data into the Sales Records Table

```
D597 Task 1=# \COPY SalesRecords(Region, Country
, ItemType, SalesChannel, OrderPriority, OrderDa
te, OrderID, ShipDate, UnitsSold, UnitPrice, Uni
tCost, TotalRevenue, TotalCost, TotalProfit) FRO
M 'C:\Users\student\Desktop\Sales Records.cs
v' DELIMITER ',' CSV HEADER;
COPY 100000
D597 Task 1=# SELECT * FROM SalesRecords LIMIT 1
0;
```

id	shipdate	region	unitsold	unitprice	unitcost	totalrevenue	totalcost	totalprofit
847	2014-10-23	Middle East and North Africa	934	152.58	97.44	142509.72	91008.96	51500.76
545	2015-02-27	Central America and the Caribbean	4551	437.20	263.33	1989697.20	1198414.83	791282.37
935	2016-01-18	Sub-Saharan Africa	9986	9.33	6.92	93169.38	69103.12	24066.26

C. Populate Products, Orders, and Reviews Tables



Tables (4)

orders

products

Columns (5)

product_id

product_name

category

sustainability

eco_impact_

Constraints

Indexes

RLS Policies

Rules

Triggers

reviews

salesrecords

Trigger Functions

Types

Views

Operators

Procedures

Sequences

Tables (4)

orders

products

reviews

salesrecords

Columns

Constraints

Indexes

RLS Policies

Rules

Triggers

Trigger Functions

Types

FTS Configurations

FTS Dictionaries

FTS Parsers

FTS Templates

Foreign Tables

Functions

Materialized Views

Operators

Procedures

Sequences

Tables (3)

orders

products

reviews

Trigger Functions

```

1 INSERT INTO Orders (Order_ID, Product_ID, Region, Units_Sold, Revenue)
2 SELECT SR.OrderID, P.Product_ID, SR.Region, SR.UnitsSOLD, SR.TotalRevenue
3 FROM SalesRecords SR
4 JOIN Products P ON SR.ItemType = P.Product_Name;
5 SELECT * FROM Orders LIMIT 10;

```

Data Output

Messages

Notifications

	order_id [PK] integer	product_id integer	region character varying (100)	units_sold integer	revenue numeric (10,2)
1	535113847	11	Middle East and North Africa	934	142509.72
2	874708545	7	Central America and the Caribbean	4551	1989697.20
3	854349935	1	Sub-Saharan Africa	9986	93169.38
4	892836844	12	Sub-Saharan Africa	9118	745214.14
5	129280602	10	Central America and the Caribbean	5858	3914725.66
6	473105037	5	Europe	1149	125562.72
7	754046475	7	Europe	7964	3481860.80
8	772153747	1	Middle East and North Africa	6307	58844.31
9	847788178	11	Europe	8217	1253749.86
10	471623599	7	Asia	2758	1205797.60

Open File

Query History

```

1 INSERT INTO Reviews (Product_ID, User_ID, Rating, Review_Text)
2 SELECT
3     P.Product_ID,
4     MOD(SR.OrderID, 1000) AS User_ID,
5     CASE
6         WHEN SR.TotalProfit > 500 THEN 5
7         WHEN SR.TotalProfit > 100 THEN 4
8         ELSE 3
9     END AS Rating,
10    'Placeholder review' AS Review_Text
11 FROM SalesRecords SR
12 JOIN Products P ON SR.ItemType = P.Product_Name;
13
14 SELECT Rating, COUNT(*) AS Count
15 FROM Reviews
16 GROUP BY Rating
17 ORDER BY Rating;

```

Data Output

Messages

Notifications

	rating integer	count bigint
1	3	42
2	4	186
3	5	99872

Query

Query History

```

1 DROP TABLE IF EXISTS SalesRecords;
2

```

Messages

Data Output

Notifications

DROP TABLE

Query returned successfully in 87 msec.

F3. Write Queries for Business Insights

Query 1: Top-Selling Products by Region

Rules

Triggers

salesrecords

Columns (14)

region

country

itemtype

saleschannel

orderpriority

orderdate

orderid

shipdate

unitsold

unitprice

unitcost

totalrevenue

totalcost

totalprofit

Constraints

Indexes

RLS Policies

Save File

Alt S

Query History

1 Region, Product_Name, SUM(Units_Sold) AS Total_Sold

2 FROM Orders

3 JOIN Products ON Orders.Product_ID = Products.Product_ID

4 GROUP BY Region, Product_Name

5 ORDER BY Total_Sold DESC;

6

Messages

Data Output

Notifications

region

product_name

total_sold

character varying (100)

character varying (255)

bigint

1 Sub-Saharan Africa

2 Sub-Saharan Africa

3 Europe

4 Sub-Saharan Africa

5 Sub-Saharan Africa

6 Sub-Saharan Africa

7 Sub-Saharan Africa

8 Europe

9 Europe

Cosmetics

Baby Food

Cereal

Beverages

Office Supplies

Household

Snacks

Office Supplies

Vegetables

11318369

11103311

11076271

11073025

11022827

10967644

10916846

10905284

10859623

Query 2: Eco-Friendly Products with High Sales

D597 Task 1/postgres@PostgreSQL 16

Query

Query History

1 SELECT Product_Name, Sustainability_Certification, Eco_Impact_Rating, SUM(Units_Sold) AS Total_Sales

2 FROM Orders

3 JOIN Products ON Orders.Product_ID = Products.Product_ID

4 WHERE Eco_Impact_Rating <= 3

5 GROUP BY Product_Name, Sustainability_Certification, Eco_Impact_Rating

6 ORDER BY Total_Sales DESC;

Data Output

Messages

Notifications

product_name

sustainability_certification

eco_impact_rating

total_sales

character varying (255)

character varying (255)

integer

bigint

1 Office Supplies

2 Cereal

3 Cosmetics

4 Baby Food

5 Clothes

6 Meat

7 Snacks

8 Personal Care

9 Beverages

Certification Placeholder

Certification Placeholder

Certification Placeholder

Certification Placeholder

Certification Placeholder

Certification Placeholder

Certification Placeholder

Certification Placeholder

Certification Placeholder

42293330

42254418

41924464

41911620

41773440

41745367

41699092

41517766

41514213

Total rows: 12 of 12

Query complete 00:00:00.234

Query 3: Top-Rated Products

F4.) Optimization of Scripts and Data Validation

a. Optimization Scripts

Pre-indexing

Query 1

Query 1

```
1 EXPLAIN ANALYZE
2 SELECT Region, Product_Name, SUM(Units_Sold) AS Total_Sold
3 FROM Orders
4 JOIN Products ON Orders.Product_ID = Products.Product_ID
5 GROUP BY Region, Product_Name
6 ORDER BY Total_Sold DESC;
```

QUERY PLAN

Step	Operation	Cost	Rows	Width	Actual Time	Actual Rows	Actual Loops
6	Batches: 1 Memory Usage: 49kB						
7	-> Hash Join (cost=11.80..2172.21 rows=100000 width=536) (actual time=0.102..40.870 rows=100000 loops=1)						
8	Hash Cond: (orders.product_id = products.product_id)						
9	-> Seq Scan on orders (cost=0.00..1884.00 rows=100000 width=24) (actual time=0.024..8.318 rows=100000 loops=1)						
10	-> Hash (cost=10.80..10.80 rows=80 width=520) (actual time=0.030..0.031 rows=12 loops=1)						
11	Buckets: 1024 Batches: 1 Memory Usage: 9kB						
12	-> Seq Scan on products (cost=0.00..10.80 rows=80 width=520) (actual time=0.016..0.019 rows=12 loops=1)						
13	Planning Time: 0.226 ms						
14	Execution Time: 81.944 ms						

Query 2

Query 2

```
1 EXPLAIN ANALYZE
2 SELECT Product_Name, Sustainability_Certification, Eco_Impact_Rating, SUM(Units_Sold) AS Total_Sales
3 FROM Orders
4 JOIN Products ON Orders.Product_ID = Products.Product_ID
5 WHERE Eco_Impact_Rating <= 3
6 GROUP BY Product_Name, Sustainability_Certification, Eco_Impact_Rating
7 ORDER BY Total_Sales DESC;
```

QUERY PLAN

Step	Operation	Cost	Rows	Width	Actual Time	Actual Rows	Actual Loops
9	-> Seq Scan on orders (cost=0.00..1884.00 rows=100000 width=24) (actual time=0.022..10.256 rows=100000 loops=1)						
10	-> Hash (cost=1.15..1.15 rows=4 width=742) (actual time=0.030..0.031 rows=12 loops=1)						
11	Buckets: 1024 Batches: 1 Memory Usage: 9kB						
12	-> Seq Scan on products (cost=0.00..1.15 rows=4 width=742) (actual time=0.016..0.021 rows=12 loops=1)						
13	Filter: (eco_impact_rating <= 3)						
14	Planning Time: 0.386 ms						
15	Execution Time: 102.533 ms						

Total rows: 15 of 15 Query complete 00:00:00.153 Ln 5, Col 29

Query 3

The screenshot shows a database management interface with a left sidebar containing a tree view of databases and schemas. The main window is divided into two tabs: 'Query' and 'Query History'. The 'Query' tab is active, displaying an SQL query. Below the query, there are tabs for 'Data Output', 'Messages', and 'Notifications'. The 'Data Output' tab is active, showing a 'QUERY PLAN' section with a table of execution details.

Query:

```
1 EXPLAIN ANALYZE
2 SELECT Product_Name, AVG(Rating) AS Avg_Rating
3 FROM Reviews
4 JOIN Products ON Reviews.Product_ID = Products.Product_ID
5 GROUP BY Product_Name
6 HAVING AVG(Rating) >= 4.5
7 ORDER BY Avg_Rating DESC;
```

QUERY PLAN

Step	Text
6	Filter: (avg(reviews.rating) >= 4.5)
7	Batches: 1 Memory Usage: 24kB
8	-> Hash Join (cost=1.27..2190.47 rows=100000 width=520) (actual time=0.131..61.98)
9	Hash Cond: (reviews.product_id = products.product_id)
10	-> Seq Scan on reviews (cost=0.00..1834.00 rows=100000 width=8) (actual time=0.00..61.98)
11	-> Hash (cost=1.12..1.12 rows=12 width=520) (actual time=0.035..0.036 rows=12)
12	Buckets: 1024 Batches: 1 Memory Usage: 9kB
13	-> Seq Scan on products (cost=0.00..1.12 rows=12 width=520) (actual time=0.00..0.03)
14	Planning Time: 0.272 ms
15	Execution Time: 110.328 ms

Indexing

The screenshot shows a database management interface with a left sidebar containing a tree view of servers, databases, and schemas. The main window is divided into two tabs: 'Query' and 'Query History'. The 'Query' tab is active, displaying an SQL query to create indexes.

Query:

```
1 CREATE INDEX idx_orders_region ON Orders (Region);
2 CREATE INDEX idx_orders_product_id ON Orders (Product_ID);
3 CREATE INDEX idx_products_eco_impact_rating ON Products (Eco_Impact_Rating);
4 CREATE INDEX idx_reviews_product_id ON Reviews (Product_ID);
5 CREATE INDEX idx_reviews_rating ON Reviews (Rating);
6
```

Post-indexing

Query 1

Query 1

```

1 EXPLAIN ANALYZE
2 SELECT Region, Product_Name, SUM(Units_Sold) AS Total_Sold
3 FROM Orders
4 JOIN Products ON Orders.Product_ID = Products.Product_ID
5 GROUP BY Region, Product_Name
6 ORDER BY Total_Sold DESC;
7
8 SET enable_seqscan = OFF;

```

Showing rows: 1 to 14 Page No: 1 of 1

Line	QUERY PLAN
5	Group Key: orders.region, products.product_name
6	Batches: 1 Memory Usage: 24kB
7	-> Hash Join (cost=12.76..5663.72 rows=100000 width=536) (actual time=0.047..43.1)
8	Hash Cond: (orders.product_id = products.product_id)
9	-> Index Scan using idx_orders_region on orders (cost=0.29..5296.05 rows=10000)
10	-> Hash (cost=12.31..12.31 rows=12 width=520) (actual time=0.016..0.018 rows=12)
11	Buckets: 1024 Batches: 1 Memory Usage: 9kB
12	-> Index Scan using products_pkey on products (cost=0.14..12.31 rows=12 width=520)
13	Planning Time: 0.245 ms
14	Execution Time: 73.183 ms

Query 2

Query 2

```

1 EXPLAIN ANALYZE
2 SELECT Product_Name, Sustainability_Certification, Eco_Impact_Rating, SUM(Units_Sold)
3 FROM Orders
4 JOIN Products ON Orders.Product_ID = Products.Product_ID
5 WHERE Eco_Impact_Rating <= 3
6 GROUP BY Product_Name, Sustainability_Certification, Eco_Impact_Rating
7 ORDER BY Total_Sales DESC;
8

```

Showing rows: 1 to 16 Page No: 1 of 1

Line	QUERY PLAN
10	-> Bitmap Heap Scan on orders (cost=92.87..749.54 rows=8333 width=8) (actual time=0.495..2.954 rows=8333)
11	Recheck Cond: (product_id = products.product_id)
12	Heap Blocks: exact=10607
13	-> Bitmap Index Scan on idx_orders_product_id (cost=0.00..90.79 rows=8333 width=0) (actual time=0.372..0.372)
14	Index Cond: (product_id = products.product_id)
15	Planning Time: 0.615 ms
16	Execution Time: 92.476 ms

Query 3

The screenshot shows a database query editor with a sidebar on the left displaying a tree view of databases and schemas. The main window is divided into two tabs: 'Query' and 'Query History'. The 'Query' tab is active, showing an SQL query:

```
1 EXPLAIN ANALYZE
2 SELECT Product_Name, AVG(Rating) AS Avg_Rating
3 FROM Reviews
4 JOIN Products ON Reviews.Product_ID = Products.Product_ID
5 GROUP BY Product_Name
6 HAVING AVG(Rating) >= 4.5
7 ORDER BY Avg_Rating DESC;
```

Below the query, there are tabs for 'Data Output', 'Messages', and 'Notifications'. The 'Data Output' tab is active, showing a 'QUERY PLAN' section with the following details:

- 6 Filter: (avg(reviews.rating) >= 4.5)
- 7 Batches: 1 Memory Usage: 24kB
- 8 -> Hash Join (cost=12.76..5536.49 rows=100000 width=520) (actual time=0.043..70.1)
- 9 Hash Cond: (reviews.product_id = products.product_id)
- 10 -> Index Scan using idx_reviews_product_id on reviews (cost=0.29..5168.82 rows=)
- 11 -> Hash (cost=12.31..12.31 rows=12 width=520) (actual time=0.015..0.017 rows=)
- 12 Buckets: 1024 Batches: 1 Memory Usage: 9kB
- 13 -> Index Scan using products_pkey on products (cost=0.14..12.31 rows=12 wic)
- 14 Planning Time: 0.227 ms
- 15 Execution Time: 105.058 ms

b. Final Data Validation

```
D597 Task 1=# \dt
          List of relations
 Schema |   Name   | Type  | Owner
-----+-----+-----+-----
 public | orders   | table | postgres
 public | products | table | postgres
 public | reviews | table | postgres
 public | salesrecords | table | postgres
(4 rows)
```

The screenshot shows a database query editor with a sidebar on the left displaying a tree view of servers and databases. The main window is divided into two tabs: 'Query' and 'Query History'. The 'Query' tab is active, showing an SQL query:

```
1 SELECT * FROM Products LIMIT 5;
```

Below the query, there are tabs for 'Data Output', 'Messages', 'Notifications', and 'Query History'. The 'Data Output' tab is active, showing a table with 5 rows and 5 columns:

	product_id [PK] integer	product_name character varying (255)	category character varying (100)	sustainability_certification character varying (100)	eco_impact_rating integer
1	1	Fruits	Category Placeholder	Certification Placeholder	3
2	2	Baby Food	Category Placeholder	Certification Placeholder	3
3	3	Beverages	Category Placeholder	Certification Placeholder	3
4	4	Vegetables	Category Placeholder	Certification Placeholder	3
5	5	Clothes	Category Placeholder	Certification Placeholder	3

Servers (1)

- PostgreSQL 16
 - Databases (2)
 - D597 Task 1
 - Casts
 - Catalogs
 - Event Triggers
 - Extensions
 - Foreign Data Wrappers
 - Languages
 - Publications
 - Schemas
 - Subscription
 - postgres
 - Casts
 - Catalogs

D597 Task 1/postgres@PostgreSQL 16

Query

```
1 SELECT * FROM Reviews LIMIT 5;
```

Data Output Messages Notifications Query History

	review_id [PK] integer	product_id integer	rating integer	review_text text
1	1	11	5	Great product!
2	2	12	5	Great product!
3	3	10	5	Great product!
4	4	2	5	Great product!
5	5	5	5	Great product!

Servers (1)

- PostgreSQL 16
 - Databases (2)
 - D597 Task 1
 - Casts
 - Catalogs
 - Event Triggers
 - Extensions
 - Foreign Data Wrappers
 - Languages
 - Publications
 - Schemas
 - Subscription
 - postgres
 - Casts
 - Catalogs

D597 Task 1/postgres@PostgreSQL 16

Query

```
1 SELECT * FROM Orders LIMIT 5;
```

Data Output Messages Notifications Query History

	order_id [PK] integer	product_id integer	region character varying (100)	units_sold integer	revenue numeric (10,2)
1	1	11	Middle East and North Africa	934	142509.72
2	2	7	Central America and the Caribbean	4551	1989697.20
3	3	1	Sub-Saharan Africa	9986	93169.38
4	4	12	Sub-Saharan Africa	9118	745214.14
5	5	10	Central America and the Caribbean	5858	3914725.66

D597 Task 1=# \di

List of relations				
Schema	Name	Type	Owner	Table
public	idx_orders_product_id	index	postgres	orders
public	idx_orders_region	index	postgres	orders
public	idx_product_id	index	postgres	orders
public	idx_product_name	index	postgres	products
public	idx_products_eco_impact_rating	index	postgres	products
public	idx_products_product_id	index	postgres	products
public	idx_region	index	postgres	orders
public	idx_reviews_product_id	index	postgres	reviews
public	idx_reviews_rating	index	postgres	reviews
public	orders_pkey	index	postgres	orders
public	products_pkey	index	postgres	products
public	reviews_pkey	index	postgres	reviews

(12 rows)

References

- Elmasri, R., & Navathe, S. B. (2016). *Fundamentals of Database Systems*. Pearson Education.
- Pang, B., & Lee, L. (2008). *Opinion Mining and Sentiment Analysis*. Foundations and Trends in Information Retrieval.
- PostgreSQL Global Development Group. (n.d.). *PostgreSQL Documentation*. Retrieved from <https://www.postgresql.org/docs/>
- Stonebraker, M., & Hellerstein, J. M. (2005). *What Goes Around Comes Around*. Communications of the ACM, 48(5), 67–72.
- United Nations. (2015). *Transforming our world: The 2030 agenda for sustainable development*. Retrieved from <https://sdgs.un.org/2030agenda>